

System project,
M1, 2011-2012

NACHOS STEP 4 : VIRTUAL MEMORY

Virtual Addresses

- The AddressSpace of a process work with virtual addresses
 - From 0 to size (size of the address space)
 - The address space of a process is also called its « virtual memory »
 - When the process is loaded into memory, the virtual addresses are to be translated into physical addresses
 - The translation is done during execution
 - When there is a memory access, the MMU translates it into the corresponding physical address
 - The translation depends on the memory management

NACHOS : MIPS Virtual Memory

- MIPS uses one of these two
 - Page table
 - TLB
- In our project, we will work only with the page table
- Do you remember what paging is and how it works?

```
class AddrSpace {
public:
    AddrSpace(OpenFile *executable);    // Create an address space,
                                        // initializing it with the p
                                        // stored in the file "executa
    ~AddrSpace();                      // De-allocate an address spac
    void InitRegisters();              // Initialize user-level CPU registers
                                        // before jumping to user code
    void SaveState();                 // Save/restore address space-specific
    void RestoreState();              // info on a context switch
private:
    TranslationEntry *pageTable;       // Assume linear page table trans
                                        // for now!
    unsigned int numPages;             // Number of pages in the virtual
                                        // address space
};
```

```
// The following class defines an entry in a translation table -- either
// in a page table or a TLB. Each entry defines a mapping from one
// virtual page to one physical page.
// In addition, there are some extra bits for access control (valid and
// read-only) and some bits for usage information (use and dirty).

class TranslationEntry {
public:
    unsigned int virtualPage; // The page number in virtual memory.
    unsigned int physicalPage; // The page number in real memory (relative to the
        // start of "mainMemory"
    bool valid; // If this bit is set, the translation is ignored.
        // (In other words, the entry hasn't been initialized.)
    bool readOnly; // If this bit is set, the user program is not allowed
        // to modify the contents of the page.
    bool use; // This bit is set by the hardware every time the
        // page is referenced or modified.
    bool dirty; // This bit is set by the hardware every time the
        // page is modified.
};
```

From virtual to physical addresses

- What is the address format?
- How do we calculate?
- Implemented in `machine/translate.cc`
 - `Translate(int virtAddr,
int* physAddr,
int size,
bool writing)`

```
//-----  
// Machine::Translate  
// Translate a virtual address into a physical address, using  
// either a page table or a TLB. Check for alignment and all sorts  
// of other errors, and if everything is ok, set the use/dirty bits in  
// the translation table entry, and store the translated physical  
// address in "physAddr". If there was an error, returns the type  
// of the exception.  
//  
// "virtAddr" -- the virtual address to translate  
// "physAddr" -- the place to store the physical address  
// "size" -- the amount of memory being read or written  
// "writing" -- if TRUE, check the "read-only" bit in the TLB  
//-----
```

ExceptionType

```
Machine::Translate(int virtAddr, int* physAddr, int size, bool writing)  
{
```

V.Marangozova, V.Danjean

Memory Access

- All accesses to the MIPS memory (main memory

```
//-----  
// Machine::ReadMem  
//     Read "size" (1, 2, or 4) bytes of virtual memory at "addr" into  
//     the location pointed to by "value".  
//  
//     Returns FALSE if the translation step from virtual to physical memor  
//     failed.  
//  
// "addr" -- the virtual address to read from  
// "size" -- the number of bytes to read (1, 2, or 4)  
// "value" -- the place to write the result  
//-----  
  
bool  
Machine::ReadMem(int addr, int size, int *value)  
{  
    int data;  
    ExceptionType exception;  
    int physicalAddress;  
  
    DEBUG('a', "Reading VA 0x%x, size %d\n", addr, size);  
}
```

ean


```
bool
Machine::ReadMem(int addr, int size, int *value)
{
    int data;
    ExceptionType exception;
    int physicalAddress;

    DEBUG('a', "Reading VA 0x%x, size %d\n", addr, size);

    exception = Translate(addr, &physicalAddress, size, FALSE);
    if (exception != NoException) {
        machine->RaiseException(exception, addr);
        return FALSE;
    }
    switch (size) {
        case 1:
            data = machine->mainMemory[physicalAddress];
            *value = data;
            break;

        case 2:
            data = *(unsigned short *) &machine->mainMemory[physicalAddress];
            *value = ShortToHost(data);
            break;

        case 4:
            data = *(unsigned int *) &machine->mainMemory[physicalAddress];
            *value = WordToHost(data);
```

For now, everything happens in physical memory directly...

- Let's go and analyse AddressSpace

```
numPages, TRUE, ,
// first, set up the translation
pageTable = new TranslationEntry[numPages];
for (i = 0; i < numPages; i++) {
    pageTable[i].virtualPage = i;    // for now, virtual page # = p
    pageTable[i].physicalPage = i;
    pageTable[i].valid = TRUE;
    pageTable[i].use = FALSE;
    pageTable[i].dirty = FALSE;
    pageTable[i].readOnly = FALSE; // if the code segment was ent
    // a separate page, we could set its
    // pages to be read-only
}

// zero out the entire address space, to zero the uninitialized data
// and the stack segment
bzero(machine->mainMemory, size);
```

```
/* then, copy in the code and data segments into memory
if (noffH.code.size > 0) {
    DEBUG('a', "Initializing code segment, at 0x%x, size %d\n",
        noffH.code.virtualAddr, noffH.code.size);
    executable->ReadAt(&(machine->mainMemory[noffH.code.virtualAddr]),
        noffH.code.size, noffH.code.inFileAddr);
}
if (noffH.initData.size > 0) {
    DEBUG('a', "Initializing data segment, at 0x%x, size %d\n",
        noffH.initData.virtualAddr, noffH.initData.size);
    executable->ReadAt(&(machine->mainMemory[noffH.initData.virtualAddr]),
        noffH.initData.size, noffH.initData.inFileAddr);
}
```

What does ReadAt?

- AddrSpace::AddrSpace in userprog/
addrspace.cc

```
//  
// "into" -- the buffer to contain the data to be read from disk  
// "from" -- the buffer containing the data to be written to disk  
// "numBytes" -- the number of bytes to transfer  
// "position" -- the offset within the file of the first byte to b  
//          read/written  
//-----  
  
int  
OpenFile::ReadAt(char *into, int numBytes, int position)  
{
```

One Subtask

- `static void ReadAtVirtual(
 OpenFile *executable,
 int virtualaddr,
 int numBytes,
 int position,
 TranslationEntry *pageTable,
 unsigned numPages)`
- Same as ReadAt but using virtual memory(pageTable and numPages).
- You could use a buffer, filled with ReadAt and then copied in memory using WriteMem

FrameProvider

- You should allocate physical pages for your processes
 - Implement FrameProvider.
 - Use de BitMap
- Bzero the memory in this class

Multiprogramming

- Syscall
 - `int ForkExec(char *s)`
 - Executable as parameter,
 - Creates a nachos thread (a main of another process),
 - in parallel with the current thread
 - Note that the new process could create threads!
- What about AddrSpace?

Points to treat...

- If there are multiple processes, the machine should not Halt at the end of the first process that terminates
- You should take care to free the memory
- mini shell